> Rami Ibrahim ING 2 TD 2
> Compte rendu TP2 BIG DATA

In this TP we'll go over Hadoop Streaming, so we need some dependencies that we need to install in order to kick start this , let's do it together !

Let's start by installing `python3` on these nodes

Let's open our terminal , move inside of the directory and execute bash on our datanode

```
(base) rami_ibrahim@fedora:~$ cd docker-hadoop/
(base) rami_ibrahim@fedora:~/docker-hadoop$ docker exec -it datanode bash
root@7b8e0c3a9630:/#
```

To install python3 we need some required reposotries, let's add them

```
bash: python: command not found
root@7b8e0c3a9630:/# > /etc/apt/sources.list
root@7b8e0c3a9630:/# echo "deb http://archive.debian.org/debian stretch main" >> /etc/apt/sour
ces.list
root@7b8e0c3a9630:/# echo "deb http://archive.debian.org/debian-security stretch/updates main"
 >> /etc/apt/sources.list
root@7b8e0c3a9630:/# apt update
Ign:1 http://archive.debian.org/debian stretch InRelease
Get:2 http://archive.debian.org/debian-security stretch/updates InRelease [59.1 kB]
Get:3 http://archive.debian.org/debian stretch Release [118 kB]
Get:4 http://archive.debian.org/debian-security stretch/updates/main amd64 Packages [782 kB]
Get:5 http://archive.debian.org/debian stretch Release.gpg [3177 B]
Get:6 http://archive.debian.org/debian stretch/main amd64 Packages [7080 kB]
Fetched 8042 kB in 1min 24s (94.7 kB/s)
Reading package lists... Done
Building dependency tree
Reading state information... Done
78 packages can be upgraded. Run 'apt list --upgradable' to see them.
```

then update the packages and run the commande to install python3

```
root@7b8e0c3a9630:/# apt install python3 python3-pip -y
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  binutils build-essential bzip2 cpp cpp-6 dbus dh-python dpkg-dev fakeroot file g++ g++-6
  gcc gcc-6 gir1.2-glib-2.0 libalgorithm-diff-perl libalgorithm-diff-xs-perl
  libalgorithm-merge-perl libapparmor1 libasan3 libatomic1 libc-dev-bin libc6-dev libcc1-0
  libcilkrts5 libdbus-1-3 libdbus-glib-1-2 libdpkg-perl libexpat1 libexpat1-dev libfakeroot
  libfile-fcntllock-perl libgcc-6-dev libgdbm3 libgirepository-1.0-1 libgomp1 libisl15
  libitm1 liblocale-gettext-perl liblsan0 libmagic-mgc libmagic1 libmpc3 libmpdec2 libmpfr4
  libmpx2 libperl5.24 libpython3-dev libpython3-stdlib libpython3.5 libpython3.5-dev
  libpython3.5-minimal libpython3.5-stdlib libquadmath0 libstdc++-6-dev libtsan0 libubsan0
  linux-libc-dev make manpages manpages-dev mime-support netbase patch perl perl-base
  perl-modules-5.24 python-pip-whl python3-cffi-backend python3-crypto python3-cryptography
```

Everything seems fine, let's check the installation

```
root@7b8e0c3a9630:/# python3 --version
Python 3.5.3
root@7b8e0c3a9630:/#
```

Once we finish the installation let's do some actual coding, we'll be needing 2 python files:
`mapper.py` and `reducer.py`

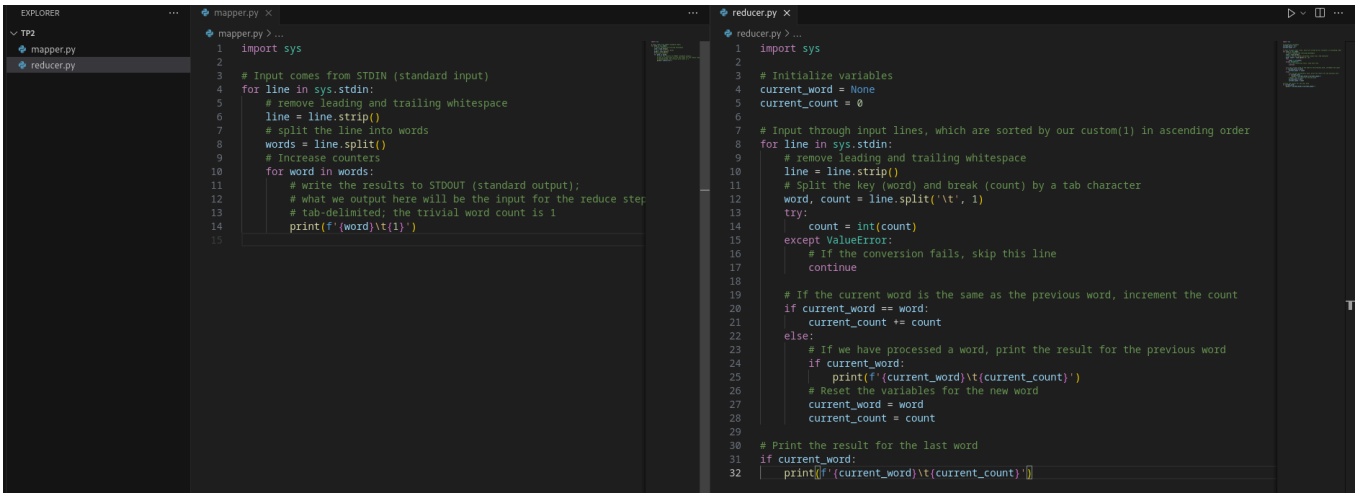

Create a text file `input.txt

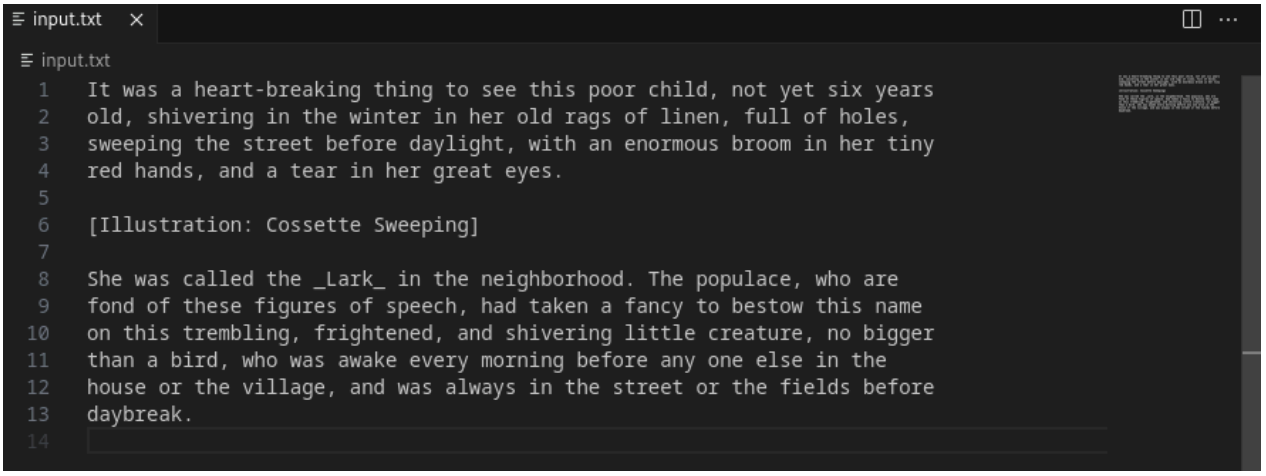

Go inside the container and create a folder named `data`

```
(base) rami_ibrahim@fedora:~/Documents/Asma/TP2$ docker exec namenode mkdir -p /data
```

Let's move these files into the container (from native)

```
(base) rami_ibrahim@fedora:~/Documents/Asma/TP2$ ls
input.txt  mapper.py  reducer.py
(base) rami_ibrahim@fedora:~/Documents/Asma/TP2$ docker cp input.txt namenode:data/input.txt
Successfully copied 2.56kB to namenode:data/input.txt
(base) rami_ibrahim@fedora:~/Documents/Asma/TP2$ docker cp mapper.py namenode:data/mapper.py
Successfully copied 2.05kB to namenode:data/mapper.py
(base) rami_ibrahim@fedora:~/Documents/Asma/TP2$ docker cp reducer.py namenode:data/reducer.py
Successfully copied 2.56kB to namenode:data/reducer.py
(base) rami_ibrahim@fedora:~/Documents/Asma/TP2$
```

Open the container bash and verify

```
(base) rami_ibrahim@fedora:~/Documents/Asma/TP2$ docker exec -it namenode bash
root@6138aafdcc59:/# cd data/
root@6138aafdcc59:/data# os
bash: os: command not found
root@6138aafdcc59:/data# ls
input.txt  mapper.py  reducer.py
root@6138aafdcc59:/data#
```

and our files are copied, let's execute them !

```
root@6138aafdcc59:/data# cat input.txt | python3 mapper.py | sort | python3 reducer.py
Cossette        1
It      1
She     1
Sweeping]       1
The     1
[Illustration:  1
_Lark_  1
a       4
always  1
an      1
and     3
any     1
are     1
awake   1
before  3
bestow  1
bigger  1
bird,   1
broom   1
called  1
child,  1
creature,       1
daybreak.       1
daylight,       1
else    1
enormous        1
every   1
eyes.   1
fancy   1
fields  1
figures 1
fond    1
frightened,     1
full    1
great   1
had     1
hands,  1
heart-breaking  1
her     3
holes,  1
house   1
in      7
linen,  1
little  1
morning 1
name    1
neighborhood.   1
no      1
not     1
of      4
old     1
old,    1
on      1
one     1
```

Now let's execute this inside of the HDFS

First, let's change the mode of the python files

```
root@6138aafdcc59:/data# chmod u+x mapper.py
root@6138aafdcc59:/data# chmod u+x reducer.py
root@6138aafdcc59:/data#
```

Now create the files and transfer them to HDFS

```
root@6138aafdcc59:/data# ls
input.txt  mapper.py  reducer.py
root@6138aafdcc59:/data# mkdir input
root@6138aafdcc59:/data# echo "hello world!">input/f1.txt
root@6138aafdcc59:/data# echo "hello docker!">input/f2.txt
root@6138aafdcc59:/data# echo "hello hadoop!">input/f3.txt
```

```
root@6138aafdcc59:/data# hdfs dfs -mkdir -p input
root@6138aafdcc59:/data# hdfs dfs -put ./input/* input
2024-12-24 18:24:22,129 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTruste
d = false, remoteHostTrusted = false
2024-12-24 18:24:22,199 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTruste
d = false, remoteHostTrusted = false
2024-12-24 18:24:22,653 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTruste
d = false, remoteHostTrusted = false
root@6138aafdcc59:/data#
```

```
root@6138aafdcc59:/data# hdfs dfs -ls
Found 1 items
drwxr-xr-x   - root supergroup          0 2024-12-24 18:24 input
root@6138aafdcc59:/data#
```

Now execute the program MapReduce

```
root@6138aafdcc59:/data# find / -name 'hadoop-streaming*.jar'
/opt/hadoop-3.2.1/share/hadoop/tools/lib/hadoop-streaming-3.2.1.jar
/opt/hadoop-3.2.1/share/hadoop/tools/sources/hadoop-streaming-3.2.1-sources.jar
/opt/hadoop-3.2.1/share/hadoop/tools/sources/hadoop-streaming-3.2.1-test-sources.jar
root@6138aafdcc59:/data#
```

```
root@e3564f682d93:/data# hadoop jar /opt/hadoop-3.2.1/share/hadoop/tools/lib/hadoop-streaming-3.2.1.jar -files mapper.py,reducer.py -input /TPs/data -output outputData
-mapper "python3 mapper.py" -reducer "python3 reducer.py"
packageJobJar: [/tmp/hadoop-unjar8093819655175065229/] [] /tmp/streamjob8967202833839866959.jar tmpDir=null
2024-12-01 21:50:04,449 INFO client.RMProxy: Connecting to ResourceManager at resourcemanager/172.18.0.3:8032
2024-12-01 21:50:04,658 INFO client.AHSProxy: Connecting to Application History server at historyserver/172.18.0.6:10200
2024-12-01 21:50:04,685 INFO client.RMProxy: Connecting to ResourceManager at resourcemanager/172.18.0.3:8032
2024-12-01 21:50:04,686 INFO client.AHSProxy: Connecting to Application History server at historyserver/172.18.0.6:10200
2024-12-01 21:50:04,870 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/root/.staging/job_1733084422932_0010
2024-12-01 21:50:05,022 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2024-12-01 21:50:05,109 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2024-12-01 21:50:05,129 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2024-12-01 21:50:05,198 INFO mapred.FileInputFormat: Total input files to process : 3
2024-12-01 21:50:05,226 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2024-12-01 21:50:05,248 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2024-12-01 21:50:05,257 INFO mapreduce.JobSubmitter: number of splits:3
2024-12-01 21:50:05,400 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2024-12-01 21:50:05,818 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1733084422932_0010
2024-12-01 21:50:05,819 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-12-01 21:50:05,968 INFO conf.Configuration: resource-types.xml not found
2024-12-01 21:50:05,969 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-12-01 21:50:06,233 INFO impl.YarnClientImpl: Submitted application application_1733084422932_0010
2024-12-01 21:50:06,265 INFO mapreduce.Job: The url to track the job: http://resourcemanager:8088/proxy/application_1733084422932_0010/
2024-12-01 21:50:06,267 INFO mapreduce.Job: Running job: job_1733084422932_0010
2024-12-01 21:50:11,345 INFO mapreduce.Job: Job job_1733084422932_0010 running in uber mode : false
2024-12-01 21:50:11,346 INFO mapreduce.Job:  map 0% reduce 0%
2024-12-01 21:50:16,404 INFO mapreduce.Job:  map 33% reduce 0%
2024-12-01 21:50:17,413 INFO mapreduce.Job:  map 67% reduce 0%
2024-12-01 21:50:18,423 INFO mapreduce.Job:  map 100% reduce 0%
2024-12-01 21:50:20,440 INFO mapreduce.Job:  map 100% reduce 100%
2024-12-01 21:50:22,467 INFO mapreduce.Job: Job job_1733084422932_0010 completed successfully
2024-12-01 21:50:22,524 INFO mapreduce.Job: Counters: 54
        File System Counters
                FILE: Number of bytes read=5128
                FILE: Number of bytes written=943932
                FILE: Number of read operations=0
                FILE: Number of large read operations=0
                FILE: Number of write operations=0
                HDFS: Number of bytes read=18103
```

And view the result

```
2024-12-01 21:50:22,524 INFO streaming.StreamJob: Output directory: outputData
root@e3564f682d93:/data# hdfs dfs -ls
Found 3 items
drwxr-xr-x   - root supergroup          0 2024-12-01 20:57 input
drwxr-xr-x   - root supergroup          0 2024-12-01 21:37 output
drwxr-xr-x   - root supergroup          0 2024-12-01 21:50 outputData
root@e3564f682d93:/data# hdfs dfs -ls outputData
Found 2 items
-rw-r--r--   3 root supergroup          0 2024-12-01 21:50 outputData/_SUCCESS
-rw-r--r--   3 root supergroup       8688 2024-12-01 21:50 outputData/part-00000
root@e3564f682d93:/data#  hdfs dfs -cat outputData/part-00000
2024-12-01 21:51:28,894 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
!important      1
"/static/assets/jupyterlab-v4/jupyterlab-index-84de2df91d1deec1912c.html";</script>     1
"all";  1
"async-google-font-2"]; 1
"{      1
&amp;   4
&rarr;</a></li> 6
'0'     1
'AIzaSyA4eNqUdRRskJsCZWVz-qL655Xa5JEMreE',      1
'Datasets'      1
'G-T7QHS60L4Q', 1
'GTM-52LNT9S',  1
'ci',   1
'content_group1':       1
'displayFeaturesTask':  1
'kaggle-161607',        1
'optimize_id':  1
'send_page_view':       1
'web-fe',       1
(a[n]=a[n]||[]).hide=h;setTimeout(function(){i();h.end=null},c);h.timeout=c;    1
(id)    1
+       1
-->     8
--><nav>        2
/>      52
/><link 3
```